

CSCI-599 Machine Learning Theory

Assignment - 3

Sree Priyanka Uppu
(uppu@usc.edu)

1

1.1 Prove that if there exists some $h \in H_k^n$ that has zero error over $S(G)$ then G is k -colorable.

Given h is consistent with $S(G)$. Let $h = \bigcap_{j=1}^k h_j$ be an ERM classifier in H_k^n over S . Using the reduction mentioned in the question, we need to prove that given k , graph G is k -colorable. Let us color the vertices of the graph as: $f(v_j) = \min\{j : h_j(e_i) = -1\}$

This algorithm, will assign a color to vertex v_i as atleast one $h_j(e_i)$ is -1 (from the reduction that, for every vertex v_i, e_i it is assigned a negative label). This implies that every vertex v_i has a color.

We have established that every vertex has a color. Next we need to show that, no 2 adjacent vertices have the same color. We will be using proof by contradiction. Assume that v_a and v_b are 2 adjacent vertices that have the same color. Then, $h_j(e_a) = h_j(e_b) = -1$. And the connecting edge will be assigned with the label: $h_j(\frac{e_a + e_b}{2})$.

Using the fact that halfspaces are convex sets and all points on a line joining a and b also lie in a convex set, implies $h_j(\frac{e_a + e_b}{2}) = -1$. This is contradictory to the definition of $S(G)$ that for every edge (v_i, v_j) construct an instance $(\frac{e_i + e_j}{2})$ with a positive label.

Thus our algorithm ensures that no 2 adjacent vertices have the same color and thus, graph G is k -colorable.

1.2 Prove that if G is k -colorable then there exists some $h \in H_k^n$ that has zero error over $S(G)$.

Given that G is k -colorable, we need to have an assignment consistent with intersection of k -half spaces.

The color assignment is, $f(v_j) = t, 1 \leq t \leq k$. Let there be k hyper planes and $h_t = \text{sgn}(w_t + b)$ where $w_{t,i} = -1$ (w_t is a normal vector to the hyper plane h_t) if $f(v_i) = t$ and $w_{t,i} = 0$ if $f(v_i) \neq t$. Given $b=0.6$ and assuming h as an intersection of hyperplanes h_1, \dots, h_k for $f(v_i) = t$ then, $h_t(E_i) = \text{sgn}(-1 + 0.6) = -1$.

As $w_{t,i} = -1$ and e_i is a unit vector in the i^{th} dimension, for the assignment of edges no 2 vertices will have the same color. Therefore, $e(v_i, v_j)$ adjacent to each other $f(v_i) = t$ and $f(v_j) = t'$ and $t \neq t'$ and we have that $w_{t,i} = 0$ and $w_{t',i} = 0$ which gives us 3 possibilities. Therefore, $h_t(\frac{e_i + e_j}{2}) = \text{sgn}(\frac{-1}{2} + 0.6) = 1$, and only will be -1 and the other vertex will be 0 at one instance of time. Hence, our algorithm is consistent with the definition of $S(G)$.

1.3 Based on the above, prove that for any $k \geq 3$, the $ERM_{H_k^n}$ problem is NP-hard.

We have shown that the reduction from k -coloring problem to the $ERM_{H_k^n}$. Also, we know that the k -coloring problem is NP-Hard and the reduction is also polynomial in time. Hence, $ERM_{H_k^n}$ is also NP-Hard.

2 AdaBoost: Show that the error of h_t w.r.t the distribution $D^{(t+1)}$ is exactly $1/2$. That is, show that for every $t \in [T]$

$$\sum_{i=1}^m D_i^{(t+1)} 1_{[y_i \neq h_t(x_i)]} = 1/2$$

Stating the AdaBoost algorithm:

1. Initialize $D_1(i) = 1/m$ for all $i=1, \dots, m$ and $T =$ number of iterations
2. for $t=1$ to T do:
 - (a). Run a weak learner L on D_t to get the hypothesis h_t which has error ϵ_t wrt D_t
 - (b). Let $\alpha_t = \frac{1}{2} \ln(\frac{1-\epsilon_t}{\epsilon_t})$
 - (c). $D_{t+1}(i) = \frac{D_t(i) \cdot \exp[-\alpha_t y_i h_t(x_i)]}{Z_t}$, where z is the normalizing factor such that $\sum_{i=1}^m D_{t+1}(i) = 1$
3. Find $H(x) = \text{sgn}(f(x))$, where $f(x) = \sum_{t=1}^T \alpha_t h_t(x)$

For a weak learner, let us assume that \exists a $h \in H$ whose $D_{t+1} < 1/2$. We already know that the normalizing factor, $Z_t = 2\sqrt{\epsilon_t(1-\epsilon_t)}$ and $\alpha_t = \frac{1}{2} \ln(\frac{1-\epsilon_t}{\epsilon_t})$, $\epsilon_t = \sum_{i: h_t(x) \neq y_i} D_t(i)$.

Since, initial distribution is uniform, $\sum_{i=1}^m D_t(i) = 1$
Solving,

$$\begin{aligned}
\sum_{i=1}^m D_i^{(t+1)} 1_{[y_i \neq h_t(x_i)]} &= \sum_{i=1}^m \frac{D_t(i) e^{-\alpha y_i h_t(x_i)} 1_{y_i h_t(x_i) < 0}}{Z_t} \\
&= \sum_{y_i h_t(x_i) < 0}^m \frac{D_t(i) e^{\alpha_t}}{Z_t} \\
&= \frac{e^{\alpha_t}}{Z_t} \sum_{y_i h_t(x_i) < 0}^m D_t(i) \\
&= \frac{\sqrt{\frac{1-\epsilon_t}{\epsilon_t}}}{2\sqrt{\epsilon_t(1-\epsilon_t)}} \epsilon_t \\
&= \frac{1}{2}
\end{aligned} \tag{1}$$

Hence proved, that the AdaBoost algorithm uses the weighting mechanism to “force” the weak learner to focus on the problematic examples in the next iteration.

3 Variant of Statistical Query Model in which the learning algorithm in addition to the Oracle $\text{STAT}(c, D)$ is also given access to the unlabeled random draws from the target distribution D .

Theorem 5.3: "Let C be a concept class and let H be a representation class over X . Then if C is efficiently learnable from statistical queries using H , C is efficiently PAC learnable using H in the presence of classification noise" holds in this variant

It implies that the learning algorithm in the noise model has access to D by ignoring the noisy labels returned by $EX_{CN}^\eta(c, D)$.

Suppose that SQ algorithm makes M queries to STAT . Use $EX_{CN}^\eta(c, D)$ to simulate each STAT query. Take enough examples each time so that the estimate \hat{p}_{chi} is off by $>\tau$ with probability $\leq \delta/M$.

We have $EX_{CN}^\eta(c, D)$ and must show how it is used efficiently to simulate $\text{STAT}(c, D)$. We have χ, τ and we need to estimate $P_\chi = \Pr_{x \sim D}[\chi(x, c(x)) = 1]$. Dividing the instance space into X_1, X_2 so that we have a part of sample where noise matters (which is X_1) and the other part being where sample doesn't matter (which is X_2). Also, the partitions are disjoint.

$X_1 = x \in X : \chi(x, 0) \neq \chi(x, 1)$ and $X_2 = x \in X : \chi(x, 0) = \chi(x, 1)$

Defining, p_1 = probability that a point x lies in X_1 . p_2 = probability that a point x lies in X_2 . $p_1 + p_2 = 1$

$p_1 = \Pr_D[X_1]$ and $p_2 = \Pr_D[X_2]$

P_χ^{CN} can be estimated from $EX_{CN}^\eta = \Pr_{EX_{CN}^\eta(c, D)}[\chi = 1]$

$P_\chi^1 = \Pr_{EX(c, D_1)}[\chi = 1]$ and $P_\chi^2 = \Pr_{EX(c, D_2)}[\chi = 1]$.

From the noisy oracle we can calculate the values of P_χ^{CN} and for p_1 since we can ignore the label for points which fall in D_1 . It will make use of the unlabeled examples available to us from D . For a large fraction of samples we can calculate p_1 using the Chernoff Bound.

Solving for,

$$\begin{aligned} P_\chi^{CN}[\chi = 1] &= (1 - \eta)P_\chi + \eta[p_1 \Pr_{EX(c, D_1)}[\chi = 0] + p_2 \Pr_{EX(c, D_2)}[\chi = 1]] \\ &= (1 - \eta)[p_1 P_\chi^1 + p_2 P_\chi^2] + \eta[p_1(1 - p_\chi^1) + p_2 P_\chi^2] \\ &= (1 - 2\eta)p_1 P_\chi^1 + p_2 P_\chi^2 + \eta p_1 \end{aligned} \tag{2}$$

$$P_\chi^1 = \frac{P_\chi^{CN} - p_2 P_\chi^2 - \eta p_1}{(1 - 2\eta)p_1} \tag{3}$$

But $P_\chi = p_1 P_\chi^1 + p_2 P_\chi^2$ Substituting in the above equation we get,

$$\frac{P_\chi - p_2 P_\chi^2}{p_1} = \frac{P_\chi^{CN} - p_2 P_\chi^2 - \eta p_1}{(1 - 2\eta)p_1} P_\chi = \frac{1}{1 - 2\eta} P_\chi^{CN} - \frac{\eta p_1}{1 - 2\eta} - \frac{2\eta}{1 - 2\eta} p_2 P_\chi^2 \tag{4}$$

We conclude that we can generate a $\text{STAT}(c, D)$ from the noisy oracle $EX_{CN}^\eta(c, D)$. In general, any PAC oracle which is using Chernoff bound to calculate the upper limit can be converted to a STAT oracle.

The concept class of axis aligned rectangles in R^n can be efficiently learned in the variant.

We use the variant of the statistical query model in which we are given access to D in addition to $\text{STAT}(c, D)$. The algorithm begins by sampling D and using the inputs drawn to partition n -dimensional space. More precisely,

for each dimension i , we use the sample to divide the x_i -axis into d/ϵ intervals with the property that the x_i component of a random point from D is approximately equally likely to fall into any of the intervals.

We now estimate the boundary of the target rectangle separately for each dimension using $\text{STAT}(c, D)$. Note that if the projection of the target rectangle onto the x_i -axis does not intersect an interval I of that axis, then the conditional probability p_I that the label is positive given that the input has its x_i component in I is 0. On the other hand, if the target's projection onto I is nonzero and there is significant probability that a positive example of the target has its x_i component in I , then p_I must be significantly larger than 0. Thus, our algorithm can start from the left, and moving to the right, place the left x_i -boundary of the hypothesis rectangle at the first interval I such that p_I is significant (at least polynomial in ϵ/n); note that estimating p_I can be done solely with calls to $\text{STAT}(c, D)$ once the intervals are defined for each coordinate. The analogous computation is done from the right, and for each dimension. The result is an efficient (polynomial in $1/\epsilon$ and n) algorithm for learning n -dimensional rectangles from statistical queries, immediately implying a noise-tolerant learning algorithm.

4 Show that if there is an efficient algorithm for PAC learning in the presence of classification noise by an algorithm that is given a noise rate upper bound η_0 ($1/2 > \eta_0 \geq \eta \geq 0$) and whose running time depends polynomially on $1/(1 - 2\eta_0)$, then there is an efficient algorithm that is given no information about the noise rate and whose running time depends polynomially on $1/(1 - 2\eta)$.

Given: an efficient PAC learnable algorithm, use the Chernoff bounds to get a bound on the number of samples required to assure that the hypothesis is correct with a probability, δ .

We have $0 \leq p \leq 1$, $0 \leq s \leq 1$ and $m > 0$ also, m is an integer. Applying

Chernoff bounds we get,

$$LE(p, m, p - s) \leq e^{-2s^2m} \quad (5)$$

$$GE(p, m, p + s) \leq e^{-2s^2m} \quad (6)$$

For a PAC learning algorithm we know the Chernoff bound to be:

$$LE(p, m, p - s) \leq \delta \quad (7)$$

$$GE(p, m, p + s) \leq \delta \quad (8)$$

We have, $e^{-2s^2m} \leq \delta$ and solving for m,

$$m \geq \frac{1}{2s^2} \ln\left(\frac{1}{\delta}\right) \quad (9)$$

Let $\sigma = (x_1, b_1), (x_2, b_2), \dots, (x_m, b_m)$ denote a series of samples from a noisy oracle $EX_\eta(c, D)$ and h is any possible hypothesis. Suppose that $F(h_i, \sigma)$ denote the number of indices, j for which h disagrees with (x_j, b_j) . Then, the rate at which a hypothesis h_i disagrees with the examples, (x_j, b_j) is equivalent to the noise rate η . Hence, to estimate the noise rate η we use the minimum disagreement over all hypotheses.

We find this estimate, η by taking samples which are sufficient such that the empirical rate of disagreement for each hypothesis is almost equal to the average. Using an iterative search procedure that tries to find a bound η_0 , by successively reducing the gap between the actual noise rate, η and $1/2$. Let us begin with an initial guess that $\eta < 1/4$ and $\eta_0 = 1/4$. During the hypothesis testing, if the value of η fails then we increase the value of $\eta = 3/8$ and then by $\eta = 7/16$ and so on. (Every time the halving distance between the previous η value and $1/2$).

Now, we can test the noise rate by:

1. Drawing some samples and estimate the failure probability of each of the hypothesis $h_i \in H$.
2. The smallest empirical failure rate, $\hat{p}_i = F(h_i, \sigma)/m$, (m = the number of samples) is compared to the current value of η_0 and if we get $\hat{p}_i < \eta_0$, we stop the search and output η_0 as the bound.
3. Else, increase η_0 and keep repeatedly increasing the size drawn at each iteration.

Say, a hypothesis h_i disagrees with a fraction p_i of samples in round r , so we get the bounds on the probability that $|\hat{p}_i - p_i| \geq 2^{-(r+2)}$. We substitute $s = 2^{-(r+2)}$, $m = m_r$ and $\delta = \delta(N2^{r+2})$ in LE, GE in the above equations:

$$LE(p, m_r, p - 2^{-(r+2)}) \leq \frac{1}{2} \frac{\delta/2}{N2^r} \quad (10)$$

$$GE(p, m_r, p + 2^{-(r+2)}) \leq \frac{1}{2} \frac{\delta/2}{N2^r} \quad (11)$$

It follows that,

$$Pr(|\hat{p}_i - p_i| \geq 2^{-(r+2)}) \geq \frac{\delta/2}{N2^r} \quad (12)$$

Summation over N hypotheses and all rounds, for any round r , we have

$$Pr(|\hat{p}_i - p_i| = 2^{-(r+2)}) \leq \delta/2 \quad (13)$$

Thus, with probability $>(1-\delta/2)$, there's a possibility of 2 events: 1. Algorithm ends on or before round, $r' = 1 + \lceil \frac{1}{\log(1-2\eta)} \rceil$ 2. When the algorithm ends, the empirical estimate of noise, $\hat{\eta}_0 > \eta$

In round r' , $\eta \leq \frac{1}{2} - \frac{1}{2^{r'}}$ and the number of samples in the round, $m_{r'}$ is sufficient to ensure that $\hat{p}_{min} \leq \eta + 2^{-(r'+2)}$ with probability $>1-\delta/2$. However,

$$\begin{aligned} \hat{\eta}_0 - 2^{-(r'+2)} &= \left(\frac{1}{2} - \frac{1}{2^{-(r'+2)}} \right) - \frac{1}{2^{(r'+2)}} \\ &> \left(\frac{1}{2} - \frac{1}{2^{r'}} \right) - \frac{1}{2^{(r'+2)}} \\ &\geq \eta + 2^{-(r'+2)} \\ &\geq \hat{p}_{min} \end{aligned} \quad (14)$$

with probability $> 1 - \delta/2$. Thus, the algorithm will end at or before round r' with the probability.

Suppose the algorithm ends at round r' . Then, $\hat{p}_{min} \geq \eta - 2^{-(r+2)}$ for m_r with probability more than $1 - \delta/2$. Since it ends, $\hat{p}_{min} \leq \hat{\eta}_0 - 2^{-(r+2)}$. Thus,

$$\eta - \frac{1}{2^{r+2}} < \hat{\eta}_0 - \frac{1}{2^{r+2}} \quad (15)$$

and hence, $\eta < \hat{\eta}_0$ with probability $> 1 - \delta/2$.

Finally, the algorithm fails when atleast one of the above conditions fails. And each possibility occurs with a probability of atleast $\delta/2$ and failure occurs with a probability of δ which is the confidence parameter from the definition of PAC learning algorithm. Assuming that the algorithm ends at round $r_0 = 1 + \lceil \frac{1}{\log(1-2\eta)} \rceil$, the total number of examples is,

$$m_{r_0} = O\left(\frac{1}{(1-2\eta)^2} \ln \left\lceil \frac{N}{(1-2\eta)\delta} \right\rceil\right) \quad (16)$$

Hence proved that can efficiently learn the concept target class in polynomial $\left(\frac{1}{(1-2\eta)}\right)$.

5 Give an efficient Statistical Query (SQ) algorithm for the class of decision lists. Analyze the complexity of your algorithm.

Let F_n be a boolean valued functions on domain $\{0,1\}^n$. Assume F_n contains the constant function 1. A probabilistic decision list c , over F_n is $(f_1, r_1), (f_2, r_2), \dots, (f_s, r_s)$ where each $f_i \in F_n$ and each $r_i \in [0, 1]$ and f_s is a constant function 1.

For any $x \in X, c(x) = r_j$ then j is the least index for $f_j(x) = 1$. The functions in F_n are tested one by one in the order of the list specified by F_n , till a f_i evaluates to 1 on which x is encountered which implies, $r_j =$ probability that $x=1$.

So, c is a probabilistic decision list with ω converging probabilities for $\omega \in [0, 1]$ if $|r_i - \omega| \geq |r_{i+1} - \omega|$ for $1 \leq i < s$.

Case (i): $\omega = 0$, then the learning algorithm learns the decision lists with decreasing probabilities, that is the lists in which $r_i \geq r_j$ for $i \leq j$.

Case (ii): $\omega = 1/2$, then the c is a decision list with decreasing probability as instances with most certain outcomes are handled at the beginning of the list, F_n .

Algorithm :

Input: $\omega \in [0, 1]$, basis $F_n = f_1, f_2, \dots, f_s$ and $\epsilon, \delta, \gamma > 0$ access to random examples of the decision lists, F_n with ω converging probability.

Output: Learning algorithm with probability of atleast $1 - \delta$ and (ϵ, γ) model of probability.

Draw a large sample, S of size m polynomial in $(1/\epsilon, s, \log 1/\delta, 1/\gamma)$ of random examples, where $s = |F_n|$

Using the obtained sample, we define $p_i =$ probability that a random positive example $(x, 1)$ is drawn given $f_i(x) = 1$. From, the ω converging decision lists, we know that $|p_1 - \omega| \geq |p_i - \omega|$ for all i . This definition helps us in identifying the first variable of the list, if our estimates \hat{p}_i are sufficiently accurate, we would expect $|\hat{p}_i - \omega|$ to be maximized when $i = 1$. the function f_i for which $|\hat{p}_i - \omega|$ is greatest is placed at the head of the hypothesis list. The remainder of the list is constructed iteratively using the part of the sample on which $f_i(x) = 0$.

Iterate through the list of s examples, if the count of $(x, b) \in S$ and $f_j(x) = 1$ is $\leq m\epsilon/4s$ then assign that j value to p and initialize \hat{p}_t to 0. Else, for the $j \in J$, \hat{p}_j is the fraction of sum of all $f_j(x) = 1$ and $b=1$ by the count of all where $f_j(x) = 1$ and choose t to be the maximum value of all $|\hat{p}_j - \omega|$

At each iteration, append L with (f_t, \hat{p}_t) and finally output L .

Analysis :

For $I \subset \{1, 2, \dots, s\}$ and $j \in \{1, 2, \dots, s\}$ let $A(I, j)$ be the set of all instances x for which $f_j(x) = 1$ and $f_i(x) = 0$ for all $i \in I$. Let

$$u(I, j) = Pr_{x \in D}[x \in A(I, j)] \text{ and } v(I, j) = Pr_{(x, b) \in EX}[b = 1 | x \in A(I, j)]$$

And let, $\hat{u}(I, j)$ and $\hat{v}(I, j)$ be the corresponding empirical estimates of these quantities.

From multiplicative form of Chernoff bounds, $u(I, j) > \epsilon/2s$

$$\hat{u}(I, j) \geq 1/2 \cdot u(I, j) \quad (17)$$

with probability $1 - \delta/(2^{s+1}s)$

Further, if $\hat{u}(I, j) > \epsilon/4s$ then the number of instances in S is atleast, $m\epsilon/4s \geq (8\epsilon^2\gamma^2) \cdot \ln(2^{s+2}s/\delta)$, by applying additive form of Chernoff bounds we have,

$$|v(I, j) - \hat{v}(I, j)| \leq \epsilon\gamma/4 \quad (18)$$

with probability $1 - \delta/(2^{s+1}s)$, assuming $\hat{u}(I, j) > \epsilon/4s$.

Thus, with probability $1 - \delta$, a sample S is chosen we have,

$$u(I, j) \leq \max(\epsilon/2s, 2\hat{u}(I, j)) \quad (19)$$

and whenever we have $\hat{u}(I, j) > \epsilon/4s$ we have,

$$|v(I, j) - \hat{v}(I, j)| < \epsilon\gamma/4 \quad (20)$$

Similarly assume, the empirical estimates $\hat{u}(I, j)$ and $\hat{v}(I, j)$ satisfy the above conditions. Now, we need to prove that this assumption implies that the algorithm's output hypothesis h is an (ϵ, γ) -good model of probability.

Suppose h is given by the list $(f_{t_1}, r'_1), \dots, (f_{t_s}, r'_s)$. Let $T_i = t_1, \dots, t_i$. To prove that h is an (ϵ, γ) good model of probability, we show that for $1 \leq i \leq s$, either

$$Pr_{x \in D}[x \in A(T_{i-1}, t_i)] \leq \epsilon/2s \quad (21)$$

or

$$Pr_{x \in D}[|h(x) - c(x)| > \gamma | x \in A(T_{i-1}, t_i)] \leq \epsilon/2 \quad (22)$$

$A(T_{i-1}, t_i)$ are disjoint which implies,

$$\begin{aligned} & Pr_{x \in D}[|h(x) - c(x)| > \gamma] \\ &= \sum_{i=1}^s Pr_{x \in D}[|h(x) - c(x)| > \gamma | x \in A(T_{i-1}, t_i)] \cdot Pr_{x \in D}[x \in A(T_{i-1}, t_i)] \quad (23) \\ &\leq \epsilon \end{aligned}$$

For the i^{th} iteration of our algorithm, the hypothesis list $(f_{t_1}, r'_1), \dots, (f_{t_{i-1}}, r'_{i-1})$. Let $C_j = A(T_{i-1}, j)$ and $p_j = v(T_{j-1}, j)$ and we observe that $\hat{p}_j = \hat{v}(T_{j-1}, j)$ (Since, all examples(x,b) in S are such that $f_k(x) = 0$ for $k \in T_{i-1}$).

If t was chosen, then $\hat{u}(T_{i-1}, t) < \epsilon/4s$, and so $\hat{u}(T_{i-1}, t) < \epsilon/2s$. Thus, $Pr_{x \in D}[x \in A(T_{i-1}, t)] \leq \epsilon/2s$ holds.

Else, that is if t was not chose, for all $j \in J$, $\hat{u}(T_{i-1}, j) > \epsilon/4s$ and thus, $|p_j - \hat{p}_j| \leq \epsilon/4$ we need to prove that, $Pr_{x \in D}[|\hat{p}_t - c(x)| > \gamma | x \in C_t] \leq \epsilon/2$.

Let u be the smallest member of J. Then $p_u = r_u$, by the definition of decision lists. Since, c is given by a list of ω converging probabilities, $|r_u - \omega| \geq |r_j - \omega|$ for $j \geq u$. Thus by our choice of t,

$$|r_j - \omega| \leq |r_u - \omega| = |p_u - \omega| \leq |\hat{p}_u - \omega| + \epsilon\gamma/4 = |\hat{p}_t - \omega| + \epsilon\gamma/4 \quad (24)$$

Suppose, $\hat{p}_t \geq \omega$ then, $r_j \leq \hat{p}_t + \epsilon\gamma/4$ for $j \in J$ and thus, $c(x) \leq \hat{p}_t + \epsilon\gamma/4 \leq \hat{p}_t + \gamma$ when $x \in C_t$. Let z be the probability that an x is chosen for which $c(x) < \hat{p}_t - \gamma$, given that $x \in C_t$:

$$z = Pr_{x \in D}[c(x) < \hat{p}_t - \gamma | x \in C_t] \quad (25)$$

$$\begin{aligned} p_t &= E_{x \in D}[c(x) | x \in C_t] \\ &\leq z(\hat{p}_t - \gamma) + (1 - z)(\hat{p}_t + \epsilon\gamma/4) \\ &\leq z(p_t + \epsilon\gamma/4 - \gamma) + (1 - z)(p_t + \epsilon\gamma/2) \\ &\leq p_t + \epsilon\gamma/2 - \gamma z \end{aligned} \quad (26)$$

This implies that, $z \leq \epsilon/2$ hence, $Pr_{x \in D}[|\hat{p}_t - c(x)| > \gamma | x \in C_t] \leq \epsilon/2$ holds.

Hence, stating the theorem, let $\omega \in [0, 1]$ be fixed and let F_n be the basis of functions. Then, the p-concept class of probabilistic decision lists over basis F_n with ω - converging probabilities is learnable with a model of probability. (assuming both ω and F_n are known). Specifically, this class can be learned in time polynomial in $(1/\epsilon, 1/\gamma, \log(1/\delta), n, |F_n|)$ and the maximum time needed to evaluate any function). Thus, also proved that the algorithm of clearly runs in polynomial time. The algorithm is learnable when noisy as well.