

CSCI-599 Machine Learning Theory

Assignment - 2

Sree Priyanka Uppu
(uppu@usc.edu)

1 PAC Learning Axis-aligned Rectangles in \mathbb{R}^n

The objective is to learn an unknown axis-aligned rectangle R in an n -dimensional space. The player receives information about R only through the following process: a random point p is chosen according to some probability distribution D . The player is given the point p together with a label indicating whether p is contained in R (a positive example) or not contained in R (a negative example).

The goal of the player is to use as few examples as possible and as little computation as possible, to pick a hypothesis rectangle R' which is a close approximation to R . The error of R' is measured as the probability that a random point is chosen from D falls in the region $R \Delta R'$, where $R \Delta R' = (R - R') \cup (R' - R)$.

We observe that the tightest fit rectangle R' is always contained in the target rectangle R . That is, $R' \subseteq R$ and hence, $R \Delta R' = R - R'$. For instance, the topmost of these strips which is shaded and denoted by T' in figure below is the region above the upper boundary of R' extended to the left and right, but below the upper boundary of R . Note that for an n dimensional rectangle there are $2n$ strips. Note that there is some overlap between the $2n$ rectangle strips at the corners. Now, if we can guarantee that the weight under D of each strip is at most $\epsilon/2n$, then we can conclude that the error of R' is at most $2n(\epsilon/2n) = \epsilon$.

Let us analyze the weight of the top strip T' . Define T to be the rectangular strip along the inside top of R which encloses exactly weight $\epsilon/2n$ under D . Clearly, T' has weight exceeding $\epsilon/2n$ under D if and only if T' includes T . Furthermore, T' includes T if and only if no point in T appears in the sample S - since if S does contain a point $p \in T$, this point has a positive label since it is contained in R , and then by definition of the tightest fit, the hypothesis rectangle R' must extend upwards into T to cover p .




Figure 1: Analysis of the error contributed by the top strip T'

By the definition of T , the probability that a single draw from the distribution D misses the region T is exactly $1 - \epsilon/2n$. The probability that m independent draws from D all miss the region T is exactly $(1 - \epsilon/2n)^m$ (Since, the probability of a conjunction of independent events is the product of the probabilities of the individual events). By Union bound, the probability that any of the $2n$ strips has weight greater than $\epsilon/2n$ is at most $2n(1 - \epsilon/2n)^m$.

Then,

$$2n(1 - \epsilon/2n)^m \leq \delta$$

$$(1 - \epsilon/2n)^m \leq \delta/2n$$

Taking natural log on both sides,

$$\ln(1 - \epsilon/2n)^m \leq \ln(\delta/2n)$$

$$m\epsilon/2n \geq -\ln(\delta/2n)$$

From the inequality, $-\ln(1 - x) > x$ Thus,

$$m \geq (2n/\epsilon)\ln(2n/\delta)$$

In summary, provided our tightest-fit algorithm takes a sample of at least $m \geq (2n/\epsilon)\ln(2n/\delta)$ examples to form its hypothesis rectangle R' , we can assert that with probability at least $1 - \delta$, R' will mis classify a new point with the probability at most ϵ .

2 Two-oracle PAC model

a. Show that if H is PAC learnable then H is PAC learnable in the two-oracle model.

Assume that C is efficiently PAC Learnable using hypothesis class, H and a learning algorithm, L . Let $h \in H$ be the hypothesis output by the learner L and $c \in C$ be the target hypothesis. The distribution D consists of examples from both the D^+ and D^- examples. Hence, $D = 1/2(D^+ + D^-)$. Choose, δ such that:

$$Pr[error_D(h) \leq \epsilon/2] \geq 1 - \delta \quad (eq.1)$$

Evaluating $error_D(h)$,

$$error_D(h) = Pr_{x \sim D}[h(x) \neq c(x)]$$

$$error_D(h) = 1/2(Pr_{x \sim D^-}[h(x) \neq c(x)] + Pr_{x \sim D^+}[h(x) \neq c(x)])$$

$$error_D(h) = 1/2(error_{D^-}(h) + error_{D^+}(h))$$

Replace the value of $error_D(h)$ in eq.1, we get

$$Pr[error_{D^-}(h) \leq \epsilon/2] \geq 1 - \delta \quad \text{and} \quad Pr[error_{D^+}(h) \leq \epsilon/2] \geq 1 - \delta$$

which implies that both of them are PAC learnable as their sample complexity is polynomial in (ϵ, δ) . Hence, H is PAC learnable in 2 oracle model with the sample complexity.

b. Show that if H is PAC learnable in the two-oracle model, then H is PAC learnable in the standard one-oracle model.

H is PAC learnable in two-oracle model, which implies that there exists a learning algorithm A such that for $c \in C$, $\epsilon > 0$ and $\delta > 0$ there exists m_- and m_+ polynomial in $1/\epsilon, 1/\delta$ and size (c) such that if we draw m_- negative examples or more and m_+ positive examples or more, with confidence $1 - \delta$, the hypothesis h output by A verifies:

$$Pr[error_{D^-}(h)] \leq \epsilon \quad \text{and} \quad Pr[error_{D^+}(h)] \leq \epsilon$$

Let D be the probability distribution over negative and positive examples. If we could draw m examples such that $m \geq \max\{m_-, m_+\}$ and m is polynomial

in $1/\epsilon, 1/\delta$ and size (c) then the 2 Oracle PAC learning implies the standard oracle PAC learning model.

$$\begin{aligned} Pr[error_D(h)] &\leq Pr[error_D(h)|c(x) = 0]Pr[c(x) = 0] + Pr[error_D(h)|c(x) = 1]Pr[c(x) = 1] \\ &\leq \epsilon(Pr[c(x) = 0] + Pr[c(x) = 1]) \leq \epsilon \end{aligned}$$

There are 2 possible cases with respect to the distribution:

Case (a.) The distribution D is not too biased that is the probability of drawing a positive examples is more than ϵ or the probability of drawing a negative examples is more than ϵ .

Let us determine the probability that we do not have enough of positive examples. Let S_m be the number of positive examples obtained when drawing m examples and the probability of drawing a positive example is ϵ . Consider the probability that we draw fewer than m positive examples in q calls to EX. This probability can be bounded by the Chernoff bounds:

$$Pr[S_m \leq (1 - \alpha)m\epsilon] \leq e^{-m\epsilon\alpha^2}$$

With $\alpha=1/2$ and $m = 2m_+/\epsilon$

$$Pr[S_m \geq m_+] \leq e^{-m_+/4}$$

And this probability this bad even occurs with probability at most $=\delta/2$. Which is

$$e^{-m_+/4} \leq \delta/2$$

Solving for m_+ yields,

$$m \geq 4\ln(2/\delta)$$

Hence, $m \geq \{\frac{2m_+}{\epsilon}, \frac{4}{\epsilon}\ln(2/\delta)\}$

The same analysis holds for bounding the probability that there are fewer than m negative examples in q calls to EX. Hence, $m \geq \{\frac{2m_-}{\epsilon}, \frac{4}{\epsilon}\ln(2/\delta)\}$. Combining we get, Hence, $m \geq \{\frac{2m_+}{\epsilon}, \frac{2m_-}{\epsilon}, \frac{8}{\epsilon}\ln(\frac{2}{\delta})\}$

Using Chernoff bounds, we could show that drawing a polynomial number of examples in $1/\epsilon, 1/\delta$ is enough to show that $m \geq \max\{m_-, m_+\}$ with higher confidence.

case(b.) D is biased towards negative or positive examples which means that $Pr[error_D(h)] \leq \epsilon$

Taken with the cases above we have shown that 2 Oracle PAC learnable is equivalent to standard one-oracle model.

3 Properties of VC Dimension

a. Monotonicity of VC Dimension, if $H' \subseteq H$ then $VCdim(H') \leq VCdim(H)$

The intuition behind this property of VC Dimension is that every hypotheses that belongs to H' are also in H , so the shattered subset of H is atleast as of H' .

Definition: Let H be a finite class. Then, clearly, for any set C we have $|H_C| \leq |H|$ and thus C cannot be shattered if $|H| < 2|C|$. This implies that $VCdim(H) \leq \log_2(|H|)$.

Applying the above definition to hypothesis class H' , $VCdim(H') \leq \log_2(|H'|)$. Applying the above definition to hypothesis class H , $VCdim(H) \leq \log_2(|H|)$.

But we know that since, $H' \subseteq H$, which implies, $|H'| \leq |H|$.

Applying log on both sides and since log is an increasing sequence, we have $\log_2(|H'|) \leq \log_2(|H|)$.

From the above definition, we get

$$VCdim(H') \leq VCdim(H)$$

b. VC Dimension versus log of class size

i. Example of a class H of functions over the real interval $X=[0,1]$ such that H is infinite while $VCdim(H)=1$

If H is a class of threshold functions where,

$$H_{Threshold} = \{h_x : x \in R\}$$

where,

$$h_x(y) = 1 \text{ if } y \leq x \text{ and } 0 \text{ otherwise}$$

The $VCdim(H_{Threshold}) = 1$ over the real interval $X=[0,1]$

ii. Example of a class H of functions over $X=[0,1]$ such that H is infinite while $VCdim(H)=\log_2(|H|)$

If H is a class of Parity functions where, For a set $I \subseteq \{1,2,...,n\}$ we define a parity function, $h_i : \{0,1\}^n \rightarrow \{0,1\}$ as follows: On a binary vector

$x=(x_1, x_2, \dots x_n) \in \{0, 1\}^n$ we have,

$$h_i(x) = \left(\sum_{i \in I} x_i \right) \mod 2$$

the VC-dimension of the class of all such parity functions, $H_{parity}^n = \{h_I : I \subset \{1, 2, \dots n\}\}$ is $\text{VCdim}(H_{parity}) = \log_2(|H|)$

c. VC Dimension of union

i. Prove that $\text{VCdim}(\cup_{i=1}^r H_i) \leq 4d \log(2d) + 2 \log(r)$

If $d = \text{VCdim}(H)$ then,

$$g_H(m) \leq \sum_{i=0}^d \binom{m}{i} = \Phi_d(m)$$

$$\begin{aligned} \Pi_{H_1 \cup H_2 \dots H_r} &\leq \Pi_{H_1} + \Pi_{H_2} + \dots + \Pi_{H_r} \\ &\leq \sum_{i=0}^d \binom{m}{i} + \sum_{i=0}^d \binom{m}{i} \dots \\ &\leq r \left(\frac{e^k}{d} \right)^d \end{aligned} \tag{1}$$

From the fact that if $m \leq d$ then $\Pi_c(m) = 2^m$

From the fact that if $m > d$ then $\Pi_c(m) \leq \left(\frac{e^k}{d} \right)^d$

$$\Pi_{H_1 \cup H_2 \dots H_r} \leq r k^d$$

Since, the union class can produce all possible labellings on the examples which at most is 2^k . Since, $\text{VCdim} < \Pi_{H_1 \cup H_2 \dots H_r}$

$$2^k < r k^d$$

Using Lemma, let $a \geq 1$ and $b < 0$ Then $x \geq 4a \log(2a) + 2b \Rightarrow x \geq a \log(x) + b$

$$r k^d \geq 4d \log(2d) + 2 \log(r)$$

Combining the equations,

$$2^k \leq 4d \log(2d) + 2 \log(r)$$

Hence, $\text{VCdim}(\cup_{i=1}^r H_i) \leq 4d \log(2d) + 2 \log(r)$

ii. Prove that for $r=2$ it holds that $VCdim(H_1 \cup H_2) \leq 2d+1$

If $d=VCdim(H)$ then,

$$g_H(m) \leq \sum_{i=0}^d \binom{m}{i} = \Phi_d(m)$$

From Sauer's Lemma,

$$\begin{aligned} \Pi_H(2d+c) &\leq \sum_{i=0}^d \binom{2d+c}{i} \\ &= 1 + \sum_{i=1}^d \binom{2d+c}{i} \\ &= 1 + \sum_{i=1}^d \binom{2d+c-1}{i} + \sum_{i=1}^d \binom{2d+c-1}{i-1} \end{aligned}$$

$$\begin{aligned} \text{Since, } \binom{n+1}{k} &= \binom{n}{k-1} + \binom{n}{k} \\ &= \sum_{i=1}^d \binom{2d+c-1}{i} + \sum_{j=0}^d \binom{2d+c-1}{j} \\ &= \sum_{i=1}^d \binom{2d+c-1}{i} + \sum_{j=d+c-1}^{2d+c-1} \binom{2d+c-1}{j} \\ &\leq \sum_{i=1}^d \binom{2d+c-1}{i} + \sum_{j=d-1}^{2d+c-1} \binom{2d+c-1}{j} \\ &= -1 + 2^{2d+c-1} \\ \Pi_{H_1 \cup H_2}(2d+c) &\leq \Pi_{H_1}(2d+1) + \Pi_{H_2}(2d+1) \\ &\leq 2^{2d+c} - 2 \\ &\leq 2^{2d+c} \end{aligned} \tag{2}$$

Thus $VCdim(H_1 \cup H_2)$ cannot be greater than $2d+c$. Here, $c=1$.

Hence, $VCdim(H_1 \cup H_2) \leq 2d+1$

4

Assume that A' be the online learning algorithm with a mistake bound of m and let S_m be the number of examples needed for the PAC learning algo-

rithm. Let $0 < \epsilon < 1$ and $0 < \delta < 1$. Let the hypothesis set be H . We propose a PAC learning algorithm A as below:

Run online algorithm A' using a new set of examples from $EX(c, D)$ at every stage. If A' 's hypothesis $h \in H$ is unchanged (no mistake) for the consecutive examples then stop and output h .

For $i=1$ to $m+1$:

- i. Get the examples from Oracle of batch size s_i .
- ii. Use the s_i examples to the online algorithm, A' sequentially
- iii. If the hypothesis h_i is wrong on an example, then stop and restart with next iteration i
- iv. If the hypothesis continues without any mistake for s_i examples the output the h_i as the solution.

If this algorithm reaches iteration $i = (M + 1)$ then it is guaranteed to stop at this iteration and the current hypothesis will have zero error.

Lets begin by defining the batch size. Assume, for the first iteration a batch size of $2^1 = 2$ examples. If the hypothesis fails for any of these consecutive examples, we increase our batch size to iteration=2 and $2^2 = 4$ examples. Generalizing the series we get the batch size to be 2^i , where i is the number of the iteration.

Now we need to prove that the assumption works for the PAC learning algorithm as well.

$$\begin{aligned}
 \Pr[PAC - Learner fails] &\leq \sum_{i=1}^{m+1} \Pr[\text{output hyp with err} > \text{in the } i^{th} \text{ iteration}] \\
 &\leq \sum_{i=1}^{m+1} \frac{\delta}{2^i} \\
 &\leq \delta \sum_{i=1}^{m+1} \frac{1}{2^i} \\
 &\leq \delta
 \end{aligned} \tag{3}$$

Hence, the maximum error is bounded still by δ .

Consider some fixed hypothesis h which the algorithm constructs. Suppose that h is bad, $\Pr[h(x)=c(x)] < (1-\epsilon)$

$\Pr[\text{the bad hypo is correct on } m_i \text{ examples in a row}] \leq (1-\epsilon)^{m_i}$ (By applying a Union Bound)

We are allowed a maximum of $\delta/2^i$ inconsistent hypothesis in each iteration. For the i^{th} iteration, if h is a bad hypothesis then we can say that the maximum probability that our hypo for the i^{th} iteration is wrong is also less than $\delta/2^i$.

So,

$$\begin{aligned}
Pr[h \text{ is wrong}] &\leq \delta/2^i \\
(1 - \epsilon)^{m_i} &\leq \delta/2^i \\
e^{-\epsilon.m_i} &\leq \delta/2^i \\
\text{Since } (1 + x) &\leq e^x \\
\text{Taking log on both sides,} & \\
-\epsilon.m_i &\leq \log(\delta/2^i) \\
\text{Solving we get,} & \\
m_i &\geq \frac{1}{\epsilon} \log \frac{2^i}{\delta}
\end{aligned} \tag{4}$$

Which gives us the minimum number of examples for each iteration. Also, we know that m is the total number of examples given to the algorithm. Hence,

$$\begin{aligned}
m &\leq \sum_{i=1}^{m+1} m_i \\
&= \sum_{i=1}^{m+1} \frac{1}{\epsilon} \log \frac{2^i}{\delta} \\
&= \sum_{i=1}^{m+1} \frac{1}{\epsilon} [i \log 2 + \log(\frac{1}{\delta})] \\
&= O(\frac{1}{\epsilon} m^2 + \frac{1}{\epsilon} m \log(\frac{1}{\delta}))
\end{aligned} \tag{5}$$

Hence, the sample complexity for our PAC algorithm is $O(\frac{1}{\epsilon} m^2 + \frac{1}{\epsilon} m \log(\frac{1}{\delta}))$.

5

Let X be the infinite set $1, 2, 3, \dots$. P_1, P_2, \dots be an infinite list of computer programs. Each program takes an input $x \in X$, runs a function $f_i : X \rightarrow \{0, 1\}$ and outputs $\{0, 1\}$. Assuming, the computer program is enumerated, such that there exists a computer program M which given an input value i , outputs P_i .

To find: a learning algorithm which is guaranteed to make $O(\log t)$ prediction mistakes where t is the smallest index such that $f=f_t$.

To find the learning algorithm which makes $O(\log t)$ mistakes, we will be using the Halving algorithm. In the current context, suppose that we (the player) have access to the predictions of N “programs.”

Algorithm:

Maintain a list of programs which have not yet made mistakes, initially including all programs.

On each trial, Make prediction based on majority vote of programs in list:

$$y' = 1 \quad \text{if } |\{i : h_i(x) = 1\}| > |\{i : h_i(x) = 0\}|, \text{ else } 0$$

Eliminate experts i in list which made mistakes $h_i(x) \neq y$

To analyze the Halving Algorithm, let W = Number of surviving programs (correct programs left in list). Initially, $W = N$. If the learner makes a mistake, then $\geq 1/2$ of the remaining experts are eliminated since the learner makes a prediction based on a majority vote of the programs. So,

$$\begin{aligned} \text{After } 1 \text{ mistake, } W &\leq 1/2(N) \\ \text{After } 2 \text{ mistakes, } W &\leq 1/2(N) \\ &\dots \\ \text{After } m \text{ mistakes, } W &\leq 2^{-m}(N) \end{aligned} \tag{6}$$

From the question, we are guaranteed that there is one function f which is perfect and will never be thrown out, so $W \geq 1$. Therefore, we have $1 \leq W \leq 2^{-m}(N)$, which implies the learner makes $m \leq \log(N)$ mistakes.

Analysis:

To bound the number of mistakes without knowing the size of the list of programs, let us assume a highly increasing function for the number of programs (or the number of examples) of the form $2^{2^{(i-1)}}$. The intuition behind the proof is: Say, if we don't find the f_t in a set of programs(or examples), then we need to look at more programs(or examples) to find the f_t .

For the increasing function $2^{2^{(i-1)}}$ which generates the number of examples. Initially, For $i=1$, there will be 2 examples and the maximum number of mistakes will be 1, since according to the halving algorithm if m is the number of examples, then there can be atmost $\log(m)$ mistakes. For $i=2$, there will be 4 examples and atmost 2 mistakes. Generalizing, assuming that the element in the sequence before the last is t , then the last element will be t^2 and the number of mistakes will be atmost $\log(t^2)=2\log(t)$.

Number of Examples	Number of mistakes
2	1
2^2	2
2^{2^2}	2^2
2^{2^3}	2^3
.	.
.	.
.	.
t^2	$2\log(t)$

To get a bound on the number of prediction mistakes, we have to sum the number of mistakes till t^2 number of examples. Note that, the number of mistakes is increasing in a geometric progression. To find the n^{th} term (which is $2\log(t)$) in a geometric progression we use the below formula,

$$t_n = t_1 \cdot r^{n-1}$$

$$2\log(t) = 1 \cdot 2^{n-1}$$

Taking log on both sides,

$$\log(2\log(t)) = n - 1$$

$$n = \log(2\log(t)) + 1$$

To find the summation of the geometric sequence,

$$S_n = a(r^n - 1)/(r - 1)$$

Here $a=1, r=2$ and substituting n ,

$$S_n = 1(2^{\log(2\log(t))+1} - 1)/(2 - 1)$$

$$S_n = (4\log(t) - 1)$$

Taking the upper bound on the sum, we have that $O(\log t)$ number of prediction mistakes. Hence, we used the halving algorithm and showed it makes $O(\log t)$ mistakes.